

# Truncated Quantile Critics (TQC)

Controlling Overestimation Bias with Distributional Reinforcement Learning

Melikşah Beşir & İlteber Konuralp

Middle East Technical University  
CENG 7822 - Reinforcement Learning

January 2026

# Outline

- 1 Introduction and Motivation
- 2 Distributional Reinforcement Learning
- 3 TQC Architecture
- 4 TQC Algorithm
- 5 TQC vs SAC Comparison
- 6 Experimental Results
- 7 Applications
- 8 Conclusion


# The Overestimation Problem in Q-Learning

**Root Cause:** The max operator in Q-learning systematically overestimates values:

$$Q(s, a) \leftarrow r + \gamma \max_{a'} Q(s', a')$$

**Why does this happen?**

- Q-values are noisy estimates (especially early in training)
- max selects the highest *estimated* value
- Noise biases selection toward overestimated actions
- Error compounds through bootstrapping



Max selects action 4 (overestimated)  
instead of optimal action 3

**Consequence:** Suboptimal policies that prefer overestimated actions

## Prior Solutions to Overestimation

### Double Q-Learning (Hasselt, 2010)

Use two Q-networks: one selects the action, the other evaluates it.

$$Q(s, a) \leftarrow r + \gamma Q_{\theta_2}(s', \arg \max_{a'} Q_{\theta_1}(s', a'))$$

*Limitation:* Still uses point estimates; can underestimate in some cases.

### Twin Critics in SAC/TD3 (Fujimoto et al., 2018; Haarnoja et al., 2018)

Maintain two critics, use the minimum for target computation:

$$y = r + \gamma \min_{i=1,2} Q_{\theta_i}(s', a')$$

*Limitation:* Arbitrary choice of minimum; ignores distribution shape.

**Can we do better by modeling the full return distribution?**

# TQC: Key Insight



$Q_1 = 5.2$   
 $Q_2 = 5.8$   
Use:  $\min = 5.2$

[4.1, 4.5, ..., 6.2, 6.8]  
Sort & truncate top 6  
Use: mean of remaining

## TQC Approach (Kuznetsov et al., 2020):

- ① Model the *distribution* of returns using quantile regression
- ② Maintain multiple critics (3), each outputting many quantiles (25)
- ③ **Truncate** the highest quantiles (most likely overestimated)
- ④ Use mean of remaining quantiles for policy optimization

# From Expected Values to Distributions

**Standard RL:** Learn expected return

$$Q(s, a) = \mathbb{E}[G_t | S_t = s, A_t = a]$$

where  $G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$


**Distributional RL:** Learn return  
*distribution*

$$Z(s, a) \sim \text{Distribution of } G_t$$

**Key Insight:**

$$Q(s, a) = \mathbb{E}[Z(s, a)]$$

The expectation is just one moment—the full distribution



**Benefits of distributions:**

- Richer gradient signal
- Risk-sensitive decisions
- Better overestimation control

## C51: Categorical Distribution (Bellemare et al., 2017)

**Idea:** Represent  $Z(s, a)$  as a categorical distribution over fixed atoms.


- Choose  $N = 51$  atoms:  $z_i \in \{V_{min}, \dots, V_{max}\}$
- Network outputs probabilities  $p_i(s, a)$  for each atom
- Distribution:  $Z_\theta(s, a) = \sum_{i=1}^N p_i \delta_{z_i}$

**Distributional Bellman Equation:**

$$Z(s, a) \stackrel{D}{=} R + \gamma Z(s', a')$$

**Training:** Project target distribution onto fixed support, minimize KL divergence.

*Limitation:* Fixed support  $[V_{min}, V_{max}]$  must be chosen a priori.



# QR-DQN: Quantile Regression (Dabney et al., 2018)

**Key Improvement:** Learn *quantile locations* instead of probabilities at fixed locations.  
**Quantile Representation:**

- Fix  $N$  quantile fractions:  $\tau_i = \frac{2i-1}{2N}$
- Network outputs quantile values  $\theta_i(s, a)$
- $\theta_i \approx F_{Z(s,a)}^{-1}(\tau_i)$  (inverse CDF)


**Quantile Regression Loss:**

$$\mathcal{L}(\theta) = \sum_{i,j} \rho_{\tau_i}(\delta_{ij})$$

where  $\rho_\tau(u) = u \cdot (\tau - \mathbb{I}[u < 0])$  is the pinball loss.

**Advantages:**

- No fixed support needed
- Adapts to any return scale



# Quantile Huber Loss

The pinball loss  $\rho_\tau(u)$  has discontinuous gradients at  $u = 0$ .


**Solution:** Use Huber smoothing (Huber, 1992) for stable training.

$$\mathcal{L}_\kappa^\tau(\delta) = |\tau - \mathbb{I}[\delta < 0]| \cdot \mathcal{L}_\kappa(\delta)$$


where the Huber loss is:

$$\mathcal{L}_\kappa(\delta) = \begin{cases} \frac{1}{2}\delta^2 & \text{if } |\delta| \leq \kappa \\ \kappa(|\delta| - \frac{1}{2}\kappa) & \text{otherwise} \end{cases}$$


Pinball ( $\tau = 0.7$ )



Huber ( $\kappa = 1$ )




Quantile Huber



Truncated Quantile Critics (Kuznetsov et al., 2020) combines:

- ① SAC framework — Actor-critic with entropy regularization
- ② Distributional critics — Quantile regression for return distributions
- ③ Truncation mechanism — Drop high quantiles to control overestimation



# Network Architecture Details


## Actor Network (same as SAC):

- Input: state  $s$
- Hidden:  $256 \times 256$  (ReLU)
- Output:  $\mu(s)$ ,  $\log \sigma(s)$
- Action:  $a = \tanh(\mu + \sigma \cdot \epsilon)$

## Critic Networks (3 independent):

- Input: state  $s$ , action  $a$
- Hidden:  $256 \times 256$  (ReLU)
- Output: 25 quantile values
- Total:  $3 \times 25 = 75$  quantiles

## Single Critic



## Key Hyperparameters:


- $N = 25$  quantiles per critic
- $N_c = 3$  critics
- $d = 2$  quantiles dropped per critic (6 total)
- Quantile fractions:  $\tau_i = \frac{2i-1}{2N} = \frac{2i-1}{50}$

# The Truncation Mechanism

**Core Idea:** High quantiles (right tail) are most susceptible to overestimation.

$$3 \times 25 = 75 \text{ quantiles}$$

(sorted across all critics)



**Why truncation works:**

- Overestimation predominantly affects upper quantiles
- Dropping  $d$  quantiles per critic removes the “optimistic tail”
- More principled than  $\min(Q_1, Q_2)$  — uses distribution shape

# Truncated Target Computation

Step-by-step target calculation:

- ① Sample next action:  $a' \sim \pi_\phi(\cdot|s')$
- ② Collect all target quantiles from all critics:

$$\mathcal{Z} = \bigcup_{i=1}^{N_c} Z_{\bar{\theta}_i}(s', a') = \{\theta_1^{(1)}, \dots, \theta_N^{(1)}, \theta_1^{(2)}, \dots, \theta_N^{(N_c)}\}$$

- ③ Sort quantiles in ascending order:

$$\mathcal{Z}_{sorted} = \text{sort}(\mathcal{Z}) = \{z_{(1)}, z_{(2)}, \dots, z_{(N \cdot N_c)}\}$$

- ④ **Truncate** — remove top  $d \cdot N_c$  quantiles:

$$\mathcal{Z}_{trunc} = \{z_{(1)}, z_{(2)}, \dots, z_{(N \cdot N_c - d \cdot N_c)}\}$$

- ⑤ Compute truncated mean for policy update:

$$\bar{Q}(s', a') = \frac{1}{|\mathcal{Z}_{trunc}|} \sum_{z \in \mathcal{Z}_{trunc}} z - \alpha \log \pi_\phi(a'|s')$$

## Critic Loss Function

Each critic  $i$  is trained to minimize the quantile Huber loss:

$$\mathcal{L}_i(\theta_i) = \frac{1}{N} \sum_{j=1}^N \frac{1}{|\mathcal{Z}_{trunc}|} \sum_{z_k \in \mathcal{Z}_{trunc}} \rho_{\tau_j}^\kappa(y_k - \theta_j^{(i)}(s, a))$$

where:

- $\theta_j^{(i)}(s, a)$  is the  $j$ -th quantile output of critic  $i$
- $y_k = r + \gamma z_k$  is the target for truncated quantile  $z_k$
- $\tau_j = \frac{2j-1}{2N}$  is the quantile fraction
- $\rho_\tau^\kappa$  is the quantile Huber loss

Intuition:

- Each quantile  $\theta_j$  is trained against ALL truncated target quantiles
- Asymmetric loss ensures  $\theta_j$  estimates the  $\tau_j$  quantile of targets
- Cross-quantile training provides rich gradient information

## Actor Loss Function

The actor is updated to maximize the truncated Q-value minus entropy cost:

$$\mathcal{L}(\phi) = \mathbb{E}_{s \sim \mathcal{D}} [\alpha \log \pi_\phi(\tilde{a}|s) - \bar{Q}(s, \tilde{a})]$$

where:

- $\tilde{a} = f_\phi(\epsilon; s)$  is a reparameterized action sample
- $\alpha$  is the temperature (learned automatically)
- $\bar{Q}$  is the truncated mean of all critics' quantiles

Computing  $\bar{Q}$  for actor update:

$$\bar{Q}(s, a) = \frac{1}{N \cdot N_c - d \cdot N_c} \sum_{i=1}^{N_c} \sum_{j=1}^{N-d} \theta_{(j)}^{(i)}(s, a)$$

Note: For actor update, quantiles are sorted *within* each critic before dropping top  $d$ .

# Complete TQC Algorithm

---

## Algorithm Truncated Quantile Critics (TQC)

---

```
1: Initialize actor  $\pi_\phi$ , critics  $\{Z_{\theta_i}\}_{i=1}^{N_c}$ , targets  $\{\bar{Z}_{\bar{\theta}_i}\}_{i=1}^{N_c}$ 
2: Initialize temperature  $\alpha$ , replay buffer  $\mathcal{D}$ 
3: for each iteration do
4:   Sample  $a \sim \pi_\phi(\cdot|s)$ , observe  $r, s'$ , store  $(s, a, r, s')$  in  $\mathcal{D}$                                 // Environment interaction
5:   for each gradient step do
6:     Sample batch from  $\mathcal{D}$                                                        // Compute truncated target
7:      $a' \sim \pi_\phi(\cdot|s')$ 
8:      $\mathcal{Z} \leftarrow \bigcup_{i=1}^{N_c} Z_{\bar{\theta}_i}(s', a')$                                      // Collect all 75 quantiles
9:      $\mathcal{Z}_{trunc} \leftarrow \text{sort}(\mathcal{Z})$ , drop top  $d \cdot N_c$                            // Keep 69 quantiles
10:     $y_k \leftarrow r + \gamma z_k$  for  $z_k \in \mathcal{Z}_{trunc}$                                          // Update critics
11:    for  $i = 1, \dots, N_c$  do
12:       $\mathcal{L}_i \leftarrow \frac{1}{N|\mathcal{Z}_{trunc}|} \sum_j \sum_k \rho_{\tau_j}^\kappa(y_k - \theta_j^{(i)})$ 
13:      Update  $\theta_i$  by gradient descent on  $\mathcal{L}_i$ 
14:    end for                                                               // Update actor and temperature
15:     $\bar{Q} \leftarrow \text{truncated mean of } \{Z_{\theta_i}(s, \tilde{a})\}$ 
16:    Update  $\phi$  by gradient ascent on  $\bar{Q}(s, \tilde{a}) - \alpha \log \pi_\phi(\tilde{a}|s)$ 
17:    Update  $\alpha$  toward target entropy
18:    Soft update:  $\bar{\theta}_i \leftarrow \tau \theta_i + (1 - \tau) \bar{\theta}_i$ 
19:  end for
20: end for
```

# Architectural Comparison

Component	SAC	TQC
Actor	Gaussian policy	Gaussian policy (same)
Critics	2 networks	3 networks
Critic output	1 scalar Q-value	25 quantile values
Total values	2	75
Target computation	$\min(Q_1, Q_2)$	Truncated mean
Overestimation control	Twin minimum	Distribution truncation
Entropy	Automatic $\alpha$	Automatic $\alpha$ (same)
Parameters	$\sim 300K$	$\sim 450K (+50\%)$
Compute per step	$1\times$	$\sim 1.5\times$

**Key Difference:** TQC trades increased computation for:

- More precise overestimation control
- Richer gradient signal from distributional learning
- Faster convergence (fewer environment steps needed)

# Overestimation Control Comparison

## SAC/TD3 Approach:

$$Q_{target} = \min(Q_1, Q_2)$$


- Binary choice between 2 values
- Can overcorrect (underestimate)
- Ignores distribution information
- Arbitrary — why exactly 2 critics?



## TQC Approach:

$$Q_{target} = \text{TruncatedMean}(\mathcal{Z})$$

- Uses 69 of 75 values
- Principled truncation of tail
- Leverages full distribution
- Tunable via  $d$  parameter



# Gradient Signal Comparison

## SAC Critic Update:

$$\mathcal{L}_{SAC} = (Q_\theta(s, a) - y)^2$$

Single gradient from single target value.

## TQC Critic Update:

$$\mathcal{L}_{TQC} = \sum_{j=1}^{25} \sum_{k=1}^{69} \rho_{\tau_j}^\kappa (y_k - \theta_j)$$

$25 \times 69 = 1725$  gradient terms per critic!

## Why More Gradients Help

- Each quantile receives feedback from entire target distribution
- Network learns both location *and* shape of return distribution
- More informative signal, especially when returns are multi-modal
- Empirically: faster convergence, particularly in early training

# MuJoCo Benchmark Results (Kuznetsov et al., 2020)

Environment	SAC	TQC	Improvement
HalfCheetah	11,520	<b>12,102</b>	+5.1%
Hopper	3,234	<b>3,512</b>	+8.6%
Walker2d	4,682	<b>5,134</b>	+9.7%
Ant	5,876	<b>6,431</b>	+9.4%
Humanoid	5,237	<b>6,048</b>	+15.5%
<b>Average</b>	6,110	<b>6,645</b>	<b>+8.8%</b>

## Key Observations:

- TQC outperforms SAC on **all** tested environments
- Largest gains on complex tasks (Humanoid: +15.5%)
- Gains come primarily from **faster convergence**, not higher asymptotic performance
- Same final performance reached in ~20% fewer environment steps

# Ablation Studies

Number of Quantiles ( $N$ ):

$N$	Avg. Score
5	5,980
10	6,312
<b>25</b>	<b>6,645</b>
50	6,598

$N = 25$  optimal: enough resolution without excessive computation.

Atoms to Drop ( $d$ ):

$d$	Avg. Score
0	6,124
1	6,489
<b>2</b>	<b>6,645</b>
3	6,587
5	6,234

$d = 2$  optimal: balances overestimation control vs. information loss.

Number of Critics ( $N_c$ ):

- $N_c = 2$ : Similar to SAC with distributional critics
- $N_c = 3$ : Best performance-compute trade-off
- $N_c = 5$ : Marginal gains, 67% more compute

# When Does TQC Help Most?

TQC provides largest benefits when:

① **High-dimensional action spaces**

- More actions  $\Rightarrow$  more overestimation opportunities
- Humanoid (17D actions): +15.5% over SAC

② **Contact-rich dynamics**

- Returns are multi-modal (success vs. failure)
- Distributional critics capture this better

③ **Limited training budget**

- TQC reaches good performance faster
- 20% sample efficiency improvement

④ **Sparse rewards / Goal-conditioned tasks**

- Combined with HER, TQC+HER often outperforms SAC+HER
- Richer gradients help in low-signal regimes

When SAC might suffice:

- Simple environments (CartPole, Pendulum)
- Very long training budgets (asymptotic performance similar)
- Compute-constrained settings

# TQC in Practice

## Robotics and Control:

- Legged locomotion (Hwangbo et al., 2019)
- Dexterous manipulation
- Autonomous driving (continuous control)

## Goal-Conditioned RL:

- TQC + HER for sparse-reward manipulation
- Navigation in complex environments
- Our project: PointMaze with TQC+HER

## Risk-Sensitive Applications:

- Distributional critics enable CVaR optimization
- Can optimize for worst-case rather than average performance
- Useful in safety-critical domains

## Available Implementations:

- sb3-contrib: from sb3\_contrib import TQC
- Clean, well-tested, drop-in replacement for SAC

# TQC Code Example

```
from sb3_contrib import TQC
from stable_baselines3.her import HerReplayBuffer

# TQC with HER for goal-conditioned learning
model = TQC(
    "MultiInputPolicy",
    env,
    learning_rate=3e-4,
    buffer_size=1_000_000,
    batch_size=256,
    tau=0.005,
    gamma=0.99,

    # TQC-specific parameters
    policy_kwargs={"n_quantiles": 25},
    top_quantiles_to_drop_per_net=2,  # d=2

    # HER integration
    replay_buffer_class=HerReplayBuffer,
    replay_buffer_kwargs={
        "n_sampled_goal": 4,
        "goal_selection_strategy": "future",
    },
)
model.learn(total_timesteps=50_000)
```

# Summary

## TQC Key Contributions:

- ➊ **Distributional Critics:** Model full return distribution via quantile regression
- ➋ **Truncation Mechanism:** Principled overestimation control by dropping high quantiles
- ➌ **Multiple Critics:** 3 critics  $\times$  25 quantiles = rich value representation

## Benefits over SAC:

- ~20% faster convergence (same final performance)
- More principled than arbitrary  $\min(Q_1, Q_2)$
- Richer gradient signal from distributional learning
- Better handling of multi-modal return distributions

## Trade-offs:

- ~50% more parameters
- ~50% more compute per update
- Additional hyperparameters ( $N, N_c, d$ )

# Key Takeaways

## For Practitioners

- Use TQC when sample efficiency matters more than compute
- Default hyperparameters ( $N = 25$ ,  $N_c = 3$ ,  $d = 2$ ) work well
- Combines naturally with HER for goal-conditioned tasks
- Available in `sb3-contrib` — easy to try!

## For Researchers

- Distributional RL provides practical benefits beyond theory
- Truncation is a simple but effective idea for overestimation control
- Open questions: adaptive truncation, learned number of quantiles

Questions?

## References |

- Bellemare, M. G., Dabney, W., and Munos, R. (2017). A distributional perspective on reinforcement learning. In *International Conference on Machine Learning*, pages 449–458.
- Dabney, W., Rowland, M., Bellemare, M. G., and Munos, R. (2018). Distributional reinforcement learning with quantile regression. In *AAAI Conference on Artificial Intelligence*.
- Fujimoto, S., Hoof, H., and Meger, D. (2018). Addressing function approximation error in actor-critic methods. In *International Conference on Machine Learning*, pages 1587–1596.
- Haarnoja, T., Zhou, A., Abbeel, P., and Levine, S. (2018). Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor. In *International Conference on Machine Learning*, pages 1861–1870.
- Hasselt, H. (2010). Double Q-learning. In *Advances in Neural Information Processing Systems*, pages 2613–2621.
- Huber, P. J. (1992). Robust estimation of a location parameter. In *Breakthroughs in Statistics*, pages 492–518.
- Hwangbo, J., Lee, J., Dosovitskiy, A., Bellicoso, D., Tsounis, V., Koltun, V., and Hutter, M. (2019). Learning agile and dynamic motor skills for legged robots. *Science Robotics*, 4(26):eaau5872.
- Kuznetsov, A., Shvechikov, P., Grishin, A., and Vetrov, D. (2020). Controlling overestimation bias with truncated mixture of continuous distributional quantile critics. In *International Conference on Machine Learning*, pages 5556–5566.

# Backup: Quantile Regression Mathematics

**Goal:** Find the  $\tau$ -quantile of a distribution.

**Definition:** The  $\tau$ -quantile  $q_\tau$  satisfies:

$$P(X \leq q_\tau) = \tau$$

**Optimization View:**  $q_\tau$  minimizes the pinball loss:

$$q_\tau = \arg \min_q \mathbb{E}[\rho_\tau(X - q)]$$

where the pinball loss is:

$$\rho_\tau(u) = u \cdot (\tau - \mathbb{I}[u < 0]) = \begin{cases} \tau \cdot u & \text{if } u \geq 0 \\ (\tau - 1) \cdot u & \text{if } u < 0 \end{cases}$$

**Intuition:**

- $\tau = 0.5$ : Symmetric loss  $\Rightarrow$  median
- $\tau > 0.5$ : Penalizes underestimates more  $\Rightarrow$  higher quantile
- $\tau < 0.5$ : Penalizes overestimates more  $\Rightarrow$  lower quantile

## Backup: Full Hyperparameter Table

Hyperparameter	Symbol	Default Value
Learning rate	$\eta$	$3 \times 10^{-4}$
Discount factor	$\gamma$	0.99
Soft update coefficient	$\tau$	0.005
Batch size	$B$	256
Replay buffer size	$ \mathcal{D} $	$10^6$
Number of critics	$N_c$	3
Quantiles per critic	$N$	25
Quantiles to drop	$d$	2
Huber threshold	$\kappa$	1.0
Initial temperature	$\alpha_0$	1.0
Target entropy	$\bar{\mathcal{H}}$	$-\dim(\mathcal{A})$
Hidden layer sizes	—	[256, 256]
Activation	—	ReLU